


Retrieval-Augmented Generation Integrated with Cloud-Based Document Archive

Adam Valík | xvalik05@stud.fit.vut.cz

Vector Databases



LLM limitations
hallucinations
cut-off date
domain-specific data
proprietary data

Fig. 1

Retrieval-Augmented Generation

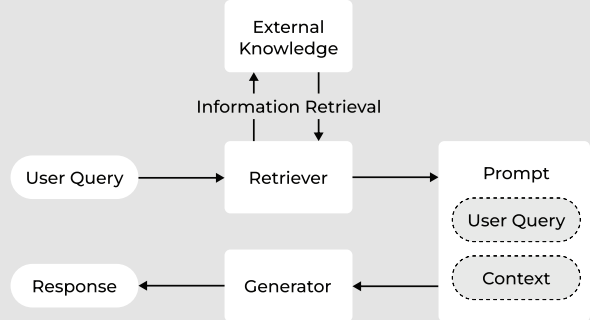


Fig. 2

Document processing

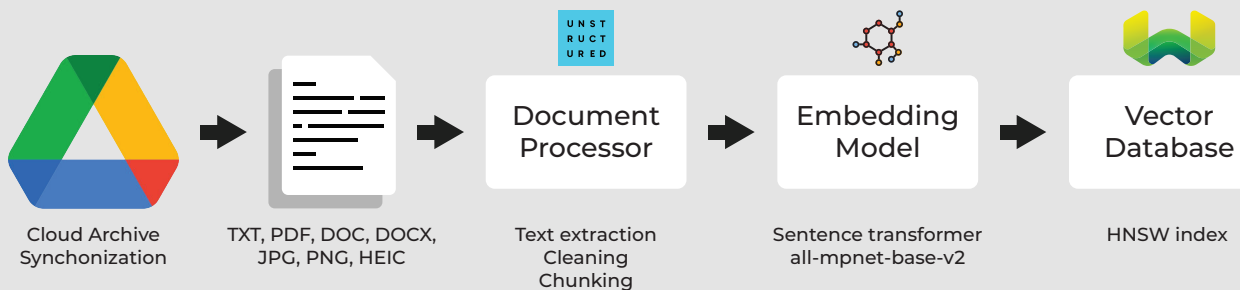


Fig. 3

Retrieval

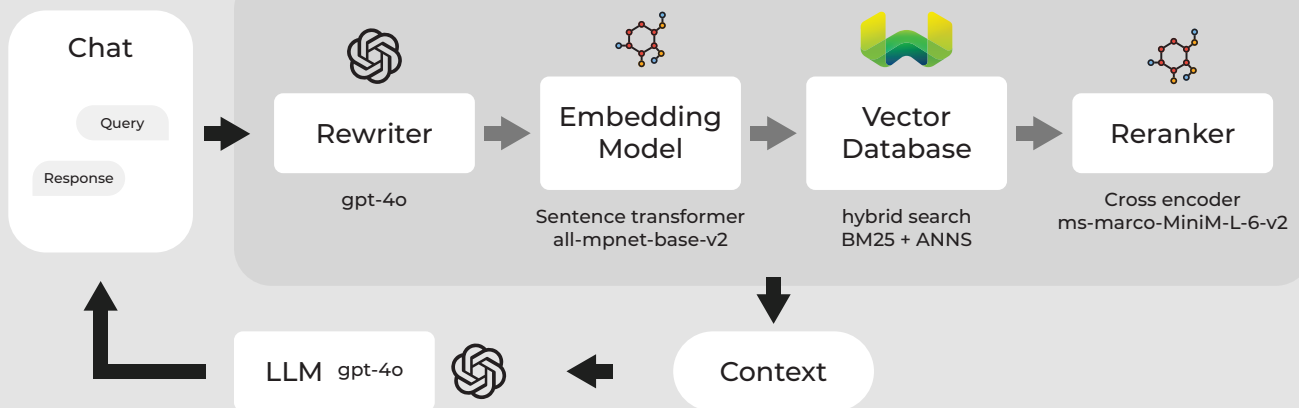


Fig. 4

Evaluation



Fig. 5

GUI

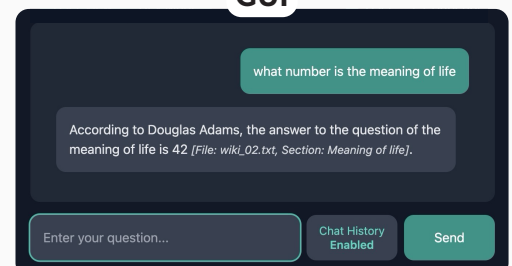


Fig. 6